

Decision Trees vs Neural Networks for Classification of Emotions from Text

Ettore Ricci

September 14, 2024

1 Introduction

In today’s digital world, the ability to interpret emotions from written text has gained importance for a variety of applications, including customer feedback analysis and mental health monitoring. Emotion classification, a challenging task within natural language processing, focuses on identifying the emotional tone within text.

This work compares the performance of three different models for emotion classification: Decision Trees, Random Forests and Neural Networks. The models were trained on a dataset of online messages labeled with 7 different emotions: Happiness, Sadness, Anger, Worry, Love, Surprise and Neutral.

Emotion	Mapping
Happiness	
Sadness	
Anger	
Worry	
Love	
Surprise	
Neutral	
Fun	Happiness
Relief	Happiness
Hate	Anger
Empty	Neutral
Enthusiasm	Happiness
Boredom	Neutral

Table 1: Emotion mappings, when the mapping is empty the emotion was not remapped.

Some instances were removed because they were duplicates. There were 23,163 duplicates and the dataset was reduced to 433,646 instances.

2 Data

The dataset used was a mix of the CrowdFlower dataset[1] and the Emotions dataset[2]. The CrowdFlower dataset contains around 40,000 tweets labeled with 13 different emotions. The Emotions dataset contains around 400,000 tweets labeled with 6 different emotions.

2.1 Exploratory data analysis and preparation

Some emotions that are present in the CrowdFlower dataset were remapped because they were not present in the Emotions dataset. The emotions and relative mappings are shown in Table 1.




Figure 1: Class distribution of the full dataset

As shown in Figure 1, the dataset is imbalanced.

The most common emotion is Happiness, and the least common is Neutral. This is because the larger dataset (Emotions) does not contain the Neutral class.




Figure 2: Text length distribution of the full dataset

The dataset was split into a training set and a test set with a 80% - 20% ratio.

2.2 Preprocessing

Multiple kinds of preprocessing were tested:

- *Bag of words (BoW)*
- *Word embeddings (WE)*
- *GloVe embeddings (GE)*

For the BoW and WE preprocessing, the text was first cleaned with a regular expression that removes urls, mentions and symbols. Then the text was tokenized and stemmed with the Snowball stemmer. Both **TFIDF** and **binary** encodings were tested for the BoW preprocessing. The GE preprocessing was done by applying the same cleaning regular expression but without stemming. The vectors used were the 6B 50d vectors from GloVe[3]. The main difference between the WE and GE preprocessing is that the WE vectors were trained on the dataset, while the GE vectors were pre-trained on a large corpus of text.

3 Models

Three different categories of models were considered:

- *Decision Trees*

- *Random Forests*
- *Neural Networks*
 - *Feedforward Networks*
 - *LSTM Networks*

Decision Trees and Random Forests both use the BoW preprocessing with **binary** encoding. Feedforward Neural Networks use the BoW preprocessing with **TFIDF** encoding. LSTM Networks were tested with both the WE and GE preprocessing.

3.1 Imbalanced Dataset

To handle the imbalanced dataset, a class weight was assigned to each class. The weight for a class c was calculated as $W_c = \frac{|D|}{|C||D_c|}$ where $|D|$ is the size of the dataset, $|C|$ is the number of classes and $|D_c|$ is the number of samples in class c . This weight was used in the criterion of the Decision Trees and Random Forests models, and in the loss function of the Neural Networks models.

3.2 Model Selection

To find a good hyperparameter configuration for each model, a random search was performed. The search was done with 10 iterations for each model (Decision Trees, Random Forests and Neural Networks). Each configuration was evaluated with a 6-fold cross-validation on the training set. Two subsequent random searches were performed, the second one narrowing the search space around the best configuration found in the first search. The search spaces for each model are shown in Table 2, Table 3 and Table 4 for the first search, and in Table 5, Table 6 and Table 7 for the second search. In the second search for the Neural Networks, early stopping was used to avoid overfitting and speed up the search.

Hyperparameter	Search Space
criterion	{gini, entropy, log_loss}
splitter	{best}
max_depth	{10, 100, None}
min_impurity_decrease	{0.01, 0.0001, 1e-06, 1e-08, 0}
class_weight	{balanced}

Table 2: Decision Trees hyperparameter search space for the first search.

Hyperparameter	Search Space
n_estimators	{10, 50, 100}
criterion	{gini, entropy, log_loss}
max_depth	{10, 100, None}
min_impurity_decrease	{0.01, 0.0001, 1e-06, 0}
n_jobs	{-1}
class_weight	{balanced}

Table 3: Random Forest hyperparameter search space for the first search.

Hyperparameter	Search Space
network	{ff_tfidf, lstm_embeddings}
base_size	{8, 16, 32}
depth	{1, 2, 3, 4}
epochs	{10, 15}
dropout	{0.5}
batchnorm	{True, False}
batch_size	{32}
lr	{0.01, 0.001}
optimizer	{adam}

Table 4: Neural Networks hyperparameter search space for the first search. While using the `ff_tfidf` network, if `batchnorm` is set to `True`, the `dropout` hyperparameter is set to 0. When using the `lstm_embeddings` and `lstm_glove` networks, the `batchnorm` hyperparameter is set to `False`.

Hyperparameter	Search Space
criterion	{gini, log_loss}
splitter	{best}
max_depth	{1000, None}
min_impurity_decrease	{1e-05, 1e-06, 1e-07}
class_weight	{balanced}

Table 5: Decision Trees hyperparameter search space for the second search.

Hyperparameter	Search Space
n_estimators	{100, 125, 150}
criterion	{gini, log_loss}
max_depth	{1000, None}
min_impurity_decrease	{1e-05, 1e-06, 1e-07}
n_jobs	{-1}
class_weight	{balanced}

Table 6: Random Forest hyperparameter search space for the second search.

Hyperparameter	Search Space
network	{ff_tfidf, lstm_glove}
base_size	{16, 32}
depth	{1, 2, 3}
epochs	{20}
patience	{2}
dropout	{0.5}
batchnorm	{False}
batch_size	{32}
lr	{0.001}
optimizer	{adam}

Table 7: Neural Networks hyperparameter search space for the second search. The same rules described in Table 4 apply.

4 Results

The models were compared using their mean accuracy in the 6-folds over the validation set. We can see the results for each model in Figure 3.




Figure 3: Accuracy distribution of each model configuration tested during the model selection process.




Figure 4: Accuracy distribution of each model category.

Among the models tested in this work, the best one was the random forest followed by the decision tree and then the neural network. The best models for each kind are shown in [Table 8](#) and their respective hyperparameters are shown in [Table 9](#), [Table 10](#) and [Table 11](#).

Model Name	Mean Accuracy
<code>decision_tree-G1-5</code>	81.02%
<code>random_forest-G1-0</code>	84.06%
<code>neural_network-G0-6</code>	66.06%

Table 8: Best models for each kind

Hyperparameter	Value
criterion	gini
splitter	best
max_depth	1000
min_impurity_decrease	1e-06
class_weight	balanced

Table 9: Hyperparameters used for model `decision_tree-G1-5`.

Hyperparameter	Value
n_estimators	125
criterion	gini
max_depth	None
min_impurity_decrease	1e-06
n_jobs	-1
class_weight	balanced

Table 10: Hyperparameters used for model `random_forest-G1-0`.

Hyperparameter	Value
network	ff_tfidf
base_size	16
depth	1
epochs	15
dropout	0.5
batchnorm	False
batch_size	32
lr	0.001
optimizer	adam

Table 11: Hyperparameters used for model `neural_network-G0-6`.

The best model was selected based on the average accuracy obtained in the cross-validation. The accuracy distribution of this model was compared to the other models with a Wilcoxon signed-rank test, some models were not significantly different ($p > 0.05$) from the best model and they are listed in Table 12.

Model	Mean Accuracy
<code>random_forest-G1-0</code>	84.06%
<code>random_forest-G1-6</code>	83.95%
<code>random_forest-G1-4</code>	83.85%

Table 12: Best model (first row) and models not significantly different from the best model.

As we can see in Figure 4, the random forest is the most susceptible to hyperparameter changes, as

it has the widest accuracy distribution, while the decision tree is the least susceptible, with the smallest accuracy distribution.

Metric	Value
Accuracy	83.98%
Precision	73.86%
Recall	83.04%
F1	76.56%

Table 13: Metrics relative to the best model `random_forest-G1-0` evaluated on the test set.




Figure 5: Confusion matrix of the best model.

The best model was also tested on a subset of the test set, composed of only the instances that were taken from the CrowdFlower dataset. The results can be compared to the results obtained by those of the paper from Batbaatar et al. [4]. The metrics obtained are shown in Table 14. We can see that the model has a worse performance on this subset compared to the model from the paper (51.1% of accuracy), this can be explained by the much more complex model used by the authors: a deep LSTM and CNN network with different kinds of embedding that extract both semantic information and sentiment information from the text.

Metric	Value
Accuracy	34.50%
Precision	33.79%
Recall	29.23%
F1	28.05%

Table 14: Metrics relative to the best model `random_forest-G1-0` evaluated on the crowdflower test set.

5 Conclusion

In this work three different kinds of machine learning algorithms were compared in the task of emotion classification from text. Random forest has proven to be the best performing model among the ones tested while the decision tree was the most stable one. The neural network was the worst performing model, but it was the one that had the most room for improvement. Unfortunately, the computational resources

available were not enough to explore more hyperparameters and configurations for the neural network.

5.1 Future Work

The neural network model has great room for improvement and maybe a more modern architecture like a transformer could have better results. Also exploring more hyperparameters for every model could lead to better results.

References

- [1] C. Van Pelt and A. Sorokin, “Designing a scalable crowdsourcing platform,” in *Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data*, pp. 765–766, 2012.
- [2] N. Elgiriye withana, “Emotions.” <https://www.kaggle.com/dsv/7563141>, 2024.
- [3] J. Pennington, R. Socher, and C. D. Manning, “Glove: Global vectors for word representation,” in *Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1532–1543, 2014.
- [4] E. Batbaatar, M. Li, and K. H. Ryu, “Semantic-emotion neural network for emotion recognition from text,” *IEEE Access*, vol. 7, pp. 111866–111878, 2019.